

Stock Market Analysis using CNN-LSTM model

This project is about analysis of Stock Market and providing suggestions and predictions to the stockholders. For this, we used CNN-LSTM approach to create a blank model, then use it to train on stock market data. Further implementation is discussed below...

In [42]:

```
# This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files under the input directory

import os
#for dirname, _, filenames in os.walk('/kaggle/input'):
#    for filename in filenames:
#        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets preserved as output when you create a version using "Save & Run All"
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside of the current session
```

Data Preprocessing and Analysis

In [43]:

```
import math
import seaborn as sns
import datetime as dt
from datetime import datetime
sns.set_style("whitegrid")
from pandas.plotting import autocorrelation_plot
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use("ggplot")
```

First we'd read the CSV file and then drop the null columns. Then we'd check the columns (some not all)

In [44]:

```
#1DP18XAREYFRWP4I
import requests
import csv
from tqdm import tqdm
key = "1DP18XAREYFRWP4I"

def request_stock_price_list(symbol, size, token):
    q_string = 'https://www.alphavantage.co/query?function=TIME_SERIES_DAILY_ADJUSTED&symbol={}&outputsize={}&apikey={}'

    print("Retrieving stock price data from Alpha Vantage (This may take a while)...")
    r = requests.get(q_string.format(symbol, size, token))
    print("Data has been successfully downloaded...")
    date = []
    colnames = list(range(0, 7))
    df = pd.DataFrame(columns = colnames)
```

```

print("Sorting the retrieved data into a dataframe...")
for i in tqdm(r.json()['Time Series (Daily)'].keys()):
    date.append(i)
    row = pd.DataFrame.from_dict(r.json()['Time Series (Daily)'][i], orient='index')
    .reset_index().T[1:]
    df = pd.concat([df, row], ignore_index=True)
df.columns = ["open", "high", "low", "close", "adjusted close", "volume", "dividend
amount", "split cf"]
df['date'] = date
return df

```

In [45]:

```

cv1 = request_stock_price_list('IBM', 'full', key)
print(cv1.head)
cv1.to_csv('data.csv')

```

Retrieving stock price data from Alpha Vantage (This may take a while)...

Data has been successfully downloaded...

Sorting the retrieved data into a dataframe...

100%|██████████| 5555/5555 [01:43<00:00, 53.74it/s]

```

<bound method NDFrame.head of
ume \
0      115.0  116.335  114.56  115.81      115.81  3322012
1      116.16  117.27  116.08  116.73      116.73  3220802
2      116.79  117.94  116.04  116.79      116.79  4914995
3      116.0   118.81  115.19  116.47      116.47  6417218
4      116.49  116.56  115.27  116.05      116.05  5384548
...
5550    92.75    92.94    90.19    90.25  52.2266076272  13737600
5551    94.44    94.44    90.0   91.56  52.9846891341  16697600
5552    95.87    95.94    93.5   94.37  54.6108029006  10369100
5553    96.75    96.81    93.69    94.81  54.8654256968  11105400
5554    98.5     98.81    96.37    96.75  55.9880807527   9551800

dividend amount split cf      date
0      0.0000      1.0  2021-11-26
1      0.0000      1.0  2021-11-24
2      0.0000      1.0  2021-11-23
3      0.0000      1.0  2021-11-22
4      0.0000      1.0  2021-11-19
...
5550    0.0000      1.0  1999-11-05
5551    0.0000      1.0  1999-11-04
5552    0.0000      1.0  1999-11-03
5553    0.0000      1.0  1999-11-02
5554    0.0000      1.0  1999-11-01

[5555 rows x 9 columns]>

```

In [46]:

```

# For data preprocessing and analysis part
data = pd.read_csv('../input/price-volume-data-for-all-us-stocks-etfs/Stocks/abe.us.txt')
#data = pd.read_csv('../input/nifty50-stock-market-data/COALINDIA.csv')
#data = pd.read_csv('../input/stock-market-data/stock_market_data/nasdaq/csv/ABCO.csv')
#data = pd.read_csv('../data.csv')
# Any CSV or TXT file can be added here....
data.dropna(inplace=True)
data.head()

```

Out[46]:

	Date	Open	High	Low	Close	Volume	OpenInt
0	2005-02-25	6.4987	6.6009	6.4668	6.5753	55766	0
1	2005-02-28	6.6072	6.7669	6.5944	6.6263	49343	0
2	2005-03-01	6.6391	6.6773	6.6072	6.6072	31643	0

3	2005-03-02	6.4987	6.6009	6.5468	6.5753	17387	0
4	2005-03-03	6.5753	6.6135	6.5562	6.5944	17387	0

In [47]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3190 entries, 0 to 3189
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        3190 non-null   object
1   Open        3190 non-null   float64
2   High        3190 non-null   float64
3   Low         3190 non-null   float64
4   Close       3190 non-null   float64
5   Volume      3190 non-null   int64
6   OpenInt     3190 non-null   int64
dtypes: float64(4), int64(2), object(1)
memory usage: 199.4+ KB
```

In [48]:

```
data.describe()
```

Out[48]:

	Open	High	Low	Close	Volume	OpenInt
count	3190.000000	3190.000000	3190.000000	3190.000000	3190.000000	3190.0
mean	11.599416	11.712848	11.484610	11.605599	28444.870846	0.0
std	2.350376	2.365621	2.327065	2.341989	37525.175821	0.0
min	5.860300	5.905000	5.834700	5.841100	106.000000	0.0
25%	10.534000	10.655000	10.413750	10.554000	8147.750000	0.0
50%	11.981000	12.067000	11.899000	11.988500	17741.500000	0.0
75%	13.271000	13.386750	13.189000	13.295750	36167.250000	0.0
max	18.130000	19.151000	17.842000	17.925000	634041.000000	0.0

In [49]:

```
data.isnull().sum()
```

Out[49]:

```
Date      0
Open      0
High      0
Low       0
Close     0
Volume    0
OpenInt   0
dtype: int64
```

In [50]:

```
data.reset_index(drop=True, inplace=True)
data.fillna(data.mean(), inplace=True)
data.head()
```

Out[50]:

	Date	Open	High	Low	Close	Volume	OpenInt
0	2005-02-25	6.4987	6.6009	6.4668	6.5753	55766	0
1	2005-02-28	6.6072	6.7660	6.5011	6.6262	102112	0

	Date	Open	High	Low	Close	Volume	OpenInt
2	2005-03-01	6.6391	6.6773	6.6072	6.6072	31643	0
3	2005-03-02	6.5753	6.6072	6.5434	6.5816	27101	0
4	2005-03-03	6.5753	6.6135	6.5562	6.5944	17387	0

In [51]:

```
data.plot(legend=True,subplots=True, figsize = (12, 6))
plt.show()
#data['Close'].plot(legend=True, figsize = (12, 6))
#plt.show()
#data['Volume'].plot(legend=True,figsize=(12,7))
#plt.show()

data.shape
data.size
data.describe(include='all').T
data.dtypes
data.nunique()
ma_day = [10,50,100]

for ma in ma_day:
    column_name = "MA for %s days" %(str(ma))
    data[column_name]=pd.DataFrame.rolling(data['Close'],ma).mean()

data['Daily Return'] = data['Close'].pct_change()
# plot the daily return percentage
data['Daily Return'].plot(figsize=(12,5),legend=True,linestyle=':',marker='o')
plt.show()

sns.displot(data['Daily Return'].dropna(),bins=100,color='green')
plt.show()

date=pd.DataFrame(data['Date'])
closing_df1 = pd.DataFrame(data['Close'])
close1 = closing_df1.rename(columns={"Close": "data_close"})
close2=pd.concat([date,close1],axis=1)
close2.head()

data.reset_index(drop=True, inplace=True)
data.fillna(data.mean(), inplace=True)
data.head()


data.nunique()




data.sort_index(axis=1,ascending=True)

cols_plot = ['Open', 'High', 'Low','Close','Volume','MA for 10 days','MA for 50 days','MA
for 100 days','Daily Return']
axes = data[cols_plot].plot(marker='.', alpha=0.7, linestyle='None', figsize=(11, 9), su
bplots=True)
for ax in axes:
    ax.set_ylabel('Daily trade')

plt.plot(data['Close'], label="Close price")
plt.xlabel("Timestamp")
plt.ylabel("Closing price")
df = data
print(df)

data.isnull().sum()
```





	Date	Open	High	Low	Close	Volume	OpenInt	\
0	2005-02-25	6.4987	6.6009	6.4668	6.5753	55766	0	
1	2005-02-28	6.6072	6.7669	6.5944	6.6263	49343	0	
2	2005-03-01	6.6391	6.6773	6.6072	6.6072	31643	0	
3	2005-03-02	6.5753	6.6072	6.5434	6.5816	27101	0	
4	2005-03-03	6.5753	6.6135	6.5562	6.5944	17387	0	
...
3185	2017-11-06	14.3998	14.4802	14.3900	14.4400	62423	0	
3186	2017-11-07	14.4400	14.4400	14.4000	14.4000	6722	0	
3187	2017-11-08	14.3400	14.4352	14.3400	14.3781	6304	0	
3188	2017-11-09	14.3300	14.3737	14.2800	14.3200	18761	0	
3189	2017-11-10	14.2500	14.3000	14.2400	14.3000	10658	0	

	MA for 10 days	MA for 50 days	MA for 100 days	Daily Return
0	11.60878	11.625272	11.649354	0.000388
1	11.60878	11.625272	11.649354	0.007756
2	11.60878	11.625272	11.649354	-0.002882
3	11.60878	11.625272	11.649354	-0.003875

4	11.60878	11.625272	11.649354	0.001945
...
3185	14.44648	14.344662	14.136796	0.003893
3186	14.43071	14.355862	14.142926	-0.002770
3187	14.42077	14.361972	14.150117	-0.001521
3188	14.40677	14.369792	14.155817	-0.004041
3189	14.39377	14.371792	14.160597	-0.001397


[3190 rows x 11 columns]

Out[51]:

```

Date          0
Open          0
High          0
Low           0
Close         0
Volume        0
OpenInt       0
MA for 10 days 0
MA for 50 days 0
MA for 100 days 0
Daily Return  0
dtype: int64

```



After that, we'll visualize the data for understanding, this is shown below...


In [52]:

```

cols_plot = ['Open', 'High', 'Low', 'Close']
axes = data[cols_plot].plot(marker='.', alpha=0.5, linestyle='None', figsize=(11, 9), su
bplots=True)
for ax in axes:
    ax.set_ylabel('Daily trade')

```





Then we'd print the data after making changes and dropping null data

In [53]:

```
plt.plot(data['Close'], label="Close price")
plt.xlabel("Timestamp")
plt.ylabel("Closing price")
df = data
print(df)


df.describe().transpose()
```

	Date	Open	High	Low	Close	Volume	OpenInt	\
0	2005-02-25	6.4987	6.6009	6.4668	6.5753	55766	0	
1	2005-02-28	6.6072	6.7669	6.5944	6.6263	49343	0	
2	2005-03-01	6.6391	6.6773	6.6072	6.6072	31643	0	
3	2005-03-02	6.5753	6.6072	6.5434	6.5816	27101	0	
4	2005-03-03	6.5753	6.6135	6.5562	6.5944	17387	0	
...	
3185	2017-11-06	14.3998	14.4802	14.3900	14.4400	62423	0	
3186	2017-11-07	14.4400	14.4400	14.4000	14.4000	6722	0	
3187	2017-11-08	14.3400	14.4352	14.3400	14.3781	6304	0	
3188	2017-11-09	14.3300	14.3737	14.2800	14.3200	18761	0	
3189	2017-11-10	14.2500	14.3000	14.2400	14.3000	10658	0	
	MA for 10 days	MA for 50 days	MA for 100 days	Daily Return				
0	11.60878	11.625272	11.649354	0.000388				
1	11.60878	11.625272	11.649354	0.007756				
2	11.60878	11.625272	11.649354	-0.002882				
3	11.60878	11.625272	11.649354	-0.003875				
4	11.60878	11.625272	11.649354	0.001945				
...				
3185	14.44648	14.344662	14.136796	0.003893				
3186	14.43071	14.355862	14.142926	-0.002770				
3187	14.42077	14.361972	14.150117	-0.001521				
3188	14.40677	14.369792	14.155817	-0.004041				
3189	14.39377	14.371792	14.160597	-0.001397				

[3190 rows x 11 columns]

Out[53]:

	count	mean	std	min	25%	50%	75%	max
Open	3190.0	11.599416	2.350376	5.860300	10.534000	11.981000	13.271000	18.130000
High	3190.0	11.712848	2.365621	5.905000	10.655000	12.067000	13.386750	19.151000
Low	3190.0	11.484610	2.327065	5.834700	10.413750	11.899000	13.189000	17.842000
Close	3190.0	11.605599	2.341989	5.841100	10.554000	11.988500	13.295750	17.925000
Volume	3190.0	28444.870846	37525.175821	106.000000	8147.750000	17741.500000	36167.250000	634041.000000
OpenInt	3190.0	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
MA for 10 days	3190.0	11.608780	2.321162	5.963080	10.577125	11.962700	13.297200	17.329400
MA for 50 days	3190.0	11.625272	2.231059	6.037646	10.591696	11.933450	13.269480	16.618460
MA for 100 days	3190.0	11.649354	2.113346	6.221377	10.632551	11.876775	13.200810	16.042560
Daily Return	3190.0	0.000388	0.017010	-0.130345	-0.006439	0.000484	0.007807	0.150503



In [54]:

```
X = data.drop(['Date', 'Close'], axis=1)
Y = data['Close']

X.shape,Y.shape

from mlxtend.feature_selection import SequentialFeatureSelector as sfs
from sklearn.linear_model import LinearRegression

lreg = LinearRegression()
sfs1 = sfs(lreg, k_features=2, forward=False, verbose=2, scoring='neg_mean_squared_error')

sfs1 = sfs1.fit(X, Y)

feat_names = list(sfs1.k_feature_names_)
print(feat_names)

# creating a new dataframe using the above variables and adding the target variable
new_data = data[feat_names]
new_data['Close'] = data['Close']

# first five rows of the new data
new_data.head()

new_data.shape, data.shape

df = new_data
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=1)]: Done 9 out of 9 | elapsed: 0.1s finished

[2021-11-29 11:49:40] Features: 8/2 -- score: -0.01130510673631423[Parallel(n_jobs=1)]: U
sing backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.0s remaining: 0.0s
[Parallel(n_jobs=1)]: Done 9 out of 9 | elapsed: 0.1s finished
```



```
[Parallel(n_jobs=1)]: Done    8 out of    8 | elapsed:    0.1s finished
```

```
[2021-11-29 11:49:40] Features: 7/2 -- score: -0.011114316874792215[Parallel(n_jobs=1)]:  
Using backend SequentialBackend with 1 concurrent workers.  
[Parallel(n_jobs=1)]: Done    1 out of    1 | elapsed:    0.0s remaining:    0.0s  
[Parallel(n_jobs=1)]: Done    7 out of    7 | elapsed:    0.1s finished
```

```
[2021-11-29 11:49:40] Features: 6/2 -- score: -0.011080371541763374[Parallel(n_jobs=1)]:  
Using backend SequentialBackend with 1 concurrent workers.  
[Parallel(n_jobs=1)]: Done    1 out of    1 | elapsed:    0.0s remaining:    0.0s  
[Parallel(n_jobs=1)]: Done    6 out of    6 | elapsed:    0.1s finished
```

```
[2021-11-29 11:49:40] Features: 5/2 -- score: -0.011080371541730121[Parallel(n_jobs=1)]:  
Using backend SequentialBackend with 1 concurrent workers.  
[Parallel(n_jobs=1)]: Done    1 out of    1 | elapsed:    0.0s remaining:    0.0s
```

```
['High', 'Low']
```

```
[Parallel(n_jobs=1)]: Done    5 out of    5 | elapsed:    0.1s finished
```

```
[2021-11-29 11:49:40] Features: 4/2 -- score: -0.011086733169734618[Parallel(n_jobs=1)]:  
Using backend SequentialBackend with 1 concurrent workers.  
[Parallel(n_jobs=1)]: Done    1 out of    1 | elapsed:    0.0s remaining:    0.0s  
[Parallel(n_jobs=1)]: Done    4 out of    4 | elapsed:    0.0s finished
```

```
[2021-11-29 11:49:40] Features: 3/2 -- score: -0.011860213917250834[Parallel(n_jobs=1)]:  
Using backend SequentialBackend with 1 concurrent workers.  
[Parallel(n_jobs=1)]: Done    1 out of    1 | elapsed:    0.0s remaining:    0.0s  
[Parallel(n_jobs=1)]: Done    3 out of    3 | elapsed:    0.0s finished
```

```
[2021-11-29 11:49:40] Features: 2/2 -- score: -0.014047232655157732/opt/conda/lib/python3  
.7/site-packages/ipykernel_launcher.py:19: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

The data has been analysed but it must be converted into data of shape [100,1] to make it easier for CNN to train on... Else it won't select necessary features and the model will fail

In [55]:

```
from sklearn.model_selection import train_test_split

X = []
Y = []
window_size=100
for i in range(1 , len(df) - window_size -1 , 1):
    first = df.iloc[i,2]
    temp = []
    temp2 = []
    for j in range(window_size):
        temp.append((df.iloc[i + j, 2] - first) / first)
    temp2.append((df.iloc[i + window_size, 2] - first) / first)
    X.append(np.array(temp).reshape(100, 1))
    Y.append(np.array(temp2).reshape(1, 1))

x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, shuffle=True)

train_X = np.array(x_train)
test_X = np.array(x_test)
train_Y = np.array(y_train)
test_Y = np.array(y_test)

train_X = train_X.reshape(train_X.shape[0],1,100,1)
test_X = test_X.reshape(test_X.shape[0],1,100,1)

print(len(train_X))
print(len(test_X))
```

Training part

This part has 2 subparts: CNN and LSTM

For CNN, the layers are created with sizes 64,128,64. In every layer, TimeDistributed function is added to track the features with respect to time. In between them, Pooling layers are added.

After that, it's passed to Bi-LSTM layers

In [56]:

```
# For creating model and training
import tensorflow as tf
from tensorflow.keras.layers import Conv1D, LSTM, Dense, Dropout, Bidirectional, TimeDistributed
from tensorflow.keras.layers import MaxPooling1D, Flatten
from tensorflow.keras.regularizers import L1, L2
from tensorflow.keras.metrics import Accuracy
from tensorflow.keras.metrics import RootMeanSquaredError

model = tf.keras.Sequential()

# Creating the Neural Network model here...
model.add(TimeDistributed(Conv1D(64, kernel_size=1, activation='relu', input_shape=(None, 100, 1))))
model.add(TimeDistributed(MaxPooling1D(2)))
model.add(TimeDistributed(Conv1D(128, kernel_size=1, activation='relu')))
model.add(TimeDistributed(MaxPooling1D(2)))
model.add(TimeDistributed(Conv1D(64, kernel_size=1, activation='relu')))
model.add(TimeDistributed(MaxPooling1D(2)))
model.add(TimeDistributed(Flatten()))
# model.add(Dense(5, kernel_regularizer=L2(0.01)))
model.add(Bidirectional(LSTM(100, return_sequences=True)))
model.add(Dropout(0.25))
model.add(Bidirectional(LSTM(100, return_sequences=False)))
model.add(Dropout(0.5))
model.add(Dense(1, activation='linear'))
model.compile(optimizer='adam', loss='mse', metrics=['mse', 'mae'])

history = model.fit(train_X, train_Y, validation_data=(test_X, test_Y), epochs=40, batch_size=40, verbose=1, shuffle=True)
```

Epoch 1/40

62/62 [=====] - 8s 35ms/step - loss: 0.0130 - mse: 0.0130 - mae: 0.0805 - val_loss: 0.0023 - val_mse: 0.0023 - val_mae: 0.0343

Epoch 2/40

62/62 [=====] - 1s 10ms/step - loss: 0.0028 - mse: 0.0028 - mae: 0.0382 - val_loss: 0.0026 - val_mse: 0.0026 - val_mae: 0.0379

Epoch 3/40

62/62 [=====] - 1s 10ms/step - loss: 0.0025 - mse: 0.0025 - mae: 0.0366 - val_loss: 0.0018 - val_mse: 0.0018 - val_mae: 0.0307

Epoch 4/40

62/62 [=====] - 1s 10ms/step - loss: 0.0022 - mse: 0.0022 - mae: 0.0341 - val_loss: 0.0021 - val_mse: 0.0021 - val_mae: 0.0349

Epoch 5/40

62/62 [=====] - 1s 10ms/step - loss: 0.0022 - mse: 0.0022 - mae: 0.0349 - val_loss: 0.0018 - val_mse: 0.0018 - val_mae: 0.0310

Epoch 6/40

62/62 [=====] - 1s 10ms/step - loss: 0.0021 - mse: 0.0021 - mae: 0.0343 - val_loss: 0.0017 - val_mse: 0.0017 - val_mae: 0.0307

Epoch 7/40

62/62 [=====] - 1s 9ms/step - loss: 0.0020 - mse: 0.0020 - mae: 0.0329 - val_loss: 0.0016 - val_mse: 0.0016 - val_mae: 0.0277

Epoch 8/40

62/62 [=====] - 1s 10ms/step - loss: 0.0021 - mse: 0.0021 - mae: 0.0328 - val_loss: 0.0014 - val_mse: 0.0014 - val_mae: 0.0265

Epoch 9/40

62/62 [=====] - 1s 10ms/step - loss: 0.0019 - mse: 0.0019 - mae: 0.0310 - val_loss: 0.0013 - val_mse: 0.0013 - val_mae: 0.0255

```
0.0315 - val_loss: 0.0016 - val_mse: 0.0016 - val_mae: 0.0282
Epoch 10/40
62/62 [=====] - 1s 10ms/step - loss: 0.0021 - mse: 0.0021 - mae:
0.0330 - val_loss: 0.0015 - val_mse: 0.0015 - val_mae: 0.0272
Epoch 11/40
62/62 [=====] - 1s 10ms/step - loss: 0.0021 - mse: 0.0021 - mae:
0.0328 - val_loss: 0.0017 - val_mse: 0.0017 - val_mae: 0.0292
Epoch 12/40
62/62 [=====] - 1s 9ms/step - loss: 0.0018 - mse: 0.0018 - mae:
0.0308 - val_loss: 0.0016 - val_mse: 0.0016 - val_mae: 0.0280
Epoch 13/40
62/62 [=====] - 1s 10ms/step - loss: 0.0019 - mse: 0.0019 - mae:
0.0321 - val_loss: 0.0016 - val_mse: 0.0016 - val_mae: 0.0281
Epoch 14/40
62/62 [=====] - 1s 11ms/step - loss: 0.0017 - mse: 0.0017 - mae:
0.0301 - val_loss: 0.0019 - val_mse: 0.0019 - val_mae: 0.0308
Epoch 15/40
62/62 [=====] - 1s 11ms/step - loss: 0.0017 - mse: 0.0017 - mae:
0.0309 - val_loss: 0.0013 - val_mse: 0.0013 - val_mae: 0.0261
Epoch 16/40
62/62 [=====] - 1s 10ms/step - loss: 0.0019 - mse: 0.0019 - mae:
0.0308 - val_loss: 0.0016 - val_mse: 0.0016 - val_mae: 0.0292
Epoch 17/40
62/62 [=====] - 1s 9ms/step - loss: 0.0017 - mse: 0.0017 - mae:
0.0303 - val_loss: 0.0013 - val_mse: 0.0013 - val_mae: 0.0260
Epoch 18/40
62/62 [=====] - 1s 10ms/step - loss: 0.0016 - mse: 0.0016 - mae:
0.0296 - val_loss: 0.0016 - val_mse: 0.0016 - val_mae: 0.0295
Epoch 19/40
62/62 [=====] - 1s 10ms/step - loss: 0.0016 - mse: 0.0016 - mae:
0.0301 - val_loss: 0.0013 - val_mse: 0.0013 - val_mae: 0.0257
Epoch 20/40
62/62 [=====] - 1s 9ms/step - loss: 0.0016 - mse: 0.0016 - mae:
0.0293 - val_loss: 0.0015 - val_mse: 0.0015 - val_mae: 0.0285
Epoch 21/40
62/62 [=====] - 1s 10ms/step - loss: 0.0017 - mse: 0.0017 - mae:
0.0300 - val_loss: 0.0017 - val_mse: 0.0017 - val_mae: 0.0288
Epoch 22/40
62/62 [=====] - 1s 10ms/step - loss: 0.0017 - mse: 0.0017 - mae:
0.0302 - val_loss: 0.0014 - val_mse: 0.0014 - val_mae: 0.0288
Epoch 23/40
62/62 [=====] - 1s 10ms/step - loss: 0.0015 - mse: 0.0015 - mae:
0.0289 - val_loss: 0.0017 - val_mse: 0.0017 - val_mae: 0.0291
Epoch 24/40
62/62 [=====] - 1s 10ms/step - loss: 0.0016 - mse: 0.0016 - mae:
0.0297 - val_loss: 0.0017 - val_mse: 0.0017 - val_mae: 0.0293
Epoch 25/40
62/62 [=====] - 1s 9ms/step - loss: 0.0016 - mse: 0.0016 - mae:
0.0296 - val_loss: 0.0018 - val_mse: 0.0018 - val_mae: 0.0307
Epoch 26/40
62/62 [=====] - 1s 10ms/step - loss: 0.0017 - mse: 0.0017 - mae:
0.0307 - val_loss: 0.0012 - val_mse: 0.0012 - val_mae: 0.0251
Epoch 27/40
62/62 [=====] - 1s 9ms/step - loss: 0.0016 - mse: 0.0016 - mae:
0.0292 - val_loss: 0.0012 - val_mse: 0.0012 - val_mae: 0.0254
Epoch 28/40
62/62 [=====] - 1s 9ms/step - loss: 0.0017 - mse: 0.0017 - mae:
0.0305 - val_loss: 0.0020 - val_mse: 0.0020 - val_mae: 0.0312
Epoch 29/40
62/62 [=====] - 1s 10ms/step - loss: 0.0019 - mse: 0.0019 - mae:
0.0312 - val_loss: 0.0014 - val_mse: 0.0014 - val_mae: 0.0266
Epoch 30/40
62/62 [=====] - 1s 10ms/step - loss: 0.0015 - mse: 0.0015 - mae:
0.0288 - val_loss: 0.0012 - val_mse: 0.0012 - val_mae: 0.0251
Epoch 31/40
62/62 [=====] - 1s 10ms/step - loss: 0.0016 - mse: 0.0016 - mae:
0.0294 - val_loss: 0.0018 - val_mse: 0.0018 - val_mae: 0.0299
Epoch 32/40
62/62 [=====] - 1s 12ms/step - loss: 0.0016 - mse: 0.0016 - mae:
0.0292 - val_loss: 0.0013 - val_mse: 0.0013 - val_mae: 0.0260
Epoch 33/40
62/62 [=====] - 1s 11ms/step - loss: 0.0015 - mse: 0.0015 - mae:
```

```

62/62 [=====] - 1s 11ms/step - loss: 0.0015 - mse: 0.0015 - mae:
0.0284 - val_loss: 0.0012 - val_mse: 0.0012 - val_mae: 0.0252
Epoch 34/40
62/62 [=====] - 1s 10ms/step - loss: 0.0015 - mse: 0.0015 - mae:
0.0290 - val_loss: 0.0012 - val_mse: 0.0012 - val_mae: 0.0243
Epoch 35/40
62/62 [=====] - 1s 10ms/step - loss: 0.0017 - mse: 0.0017 - mae:
0.0308 - val_loss: 0.0013 - val_mse: 0.0013 - val_mae: 0.0259
Epoch 36/40
62/62 [=====] - 1s 10ms/step - loss: 0.0014 - mse: 0.0014 - mae:
0.0284 - val_loss: 0.0012 - val_mse: 0.0012 - val_mae: 0.0252
Epoch 37/40
62/62 [=====] - 1s 10ms/step - loss: 0.0015 - mse: 0.0015 - mae:
0.0289 - val_loss: 0.0012 - val_mse: 0.0012 - val_mae: 0.0256
Epoch 38/40
62/62 [=====] - 1s 10ms/step - loss: 0.0015 - mse: 0.0015 - mae:
0.0290 - val_loss: 0.0013 - val_mse: 0.0013 - val_mae: 0.0265
Epoch 39/40
62/62 [=====] - 1s 9ms/step - loss: 0.0015 - mse: 0.0015 - mae:
0.0285 - val_loss: 0.0012 - val_mse: 0.0012 - val_mae: 0.0253
Epoch 40/40
62/62 [=====] - 1s 10ms/step - loss: 0.0014 - mse: 0.0014 - mae:
0.0276 - val_loss: 0.0015 - val_mse: 0.0015 - val_mae: 0.0277

```

In [57]:


```

plt.plot(history.history['loss'], label='train loss')
plt.plot(history.history['val_loss'], label='val loss')
plt.xlabel("epoch")
plt.ylabel("Loss")
plt.legend()

```

Out[57]:

<matplotlib.legend.Legend at 0x7f382ca22110>



In [58]:


```


plt.plot(history.history['mse'], label='train mse')
plt.plot(history.history['val_mse'], label='val mse')
plt.xlabel("epoch")
plt.ylabel("Loss")
plt.legend()

```

Out[58]:

<matplotlib.legend.Legend at 0x7f382c421790>






In [59]:

```
plt.plot(history.history['mae'], label='train mae')
plt.plot(history.history['val_mae'], label='val mae')
plt.xlabel("epoch")
plt.ylabel("Loss")
plt.legend()
```

Out[59]:

<matplotlib.legend.Legend at 0x7f3335291410>



In [60]:

```
# After the model has been constructed, we need to train
from tensorflow.keras.utils import plot_model
print(model.summary())
plot_model(model, to_file='model.png', show_shapes=True, show_layer_names=True)
```

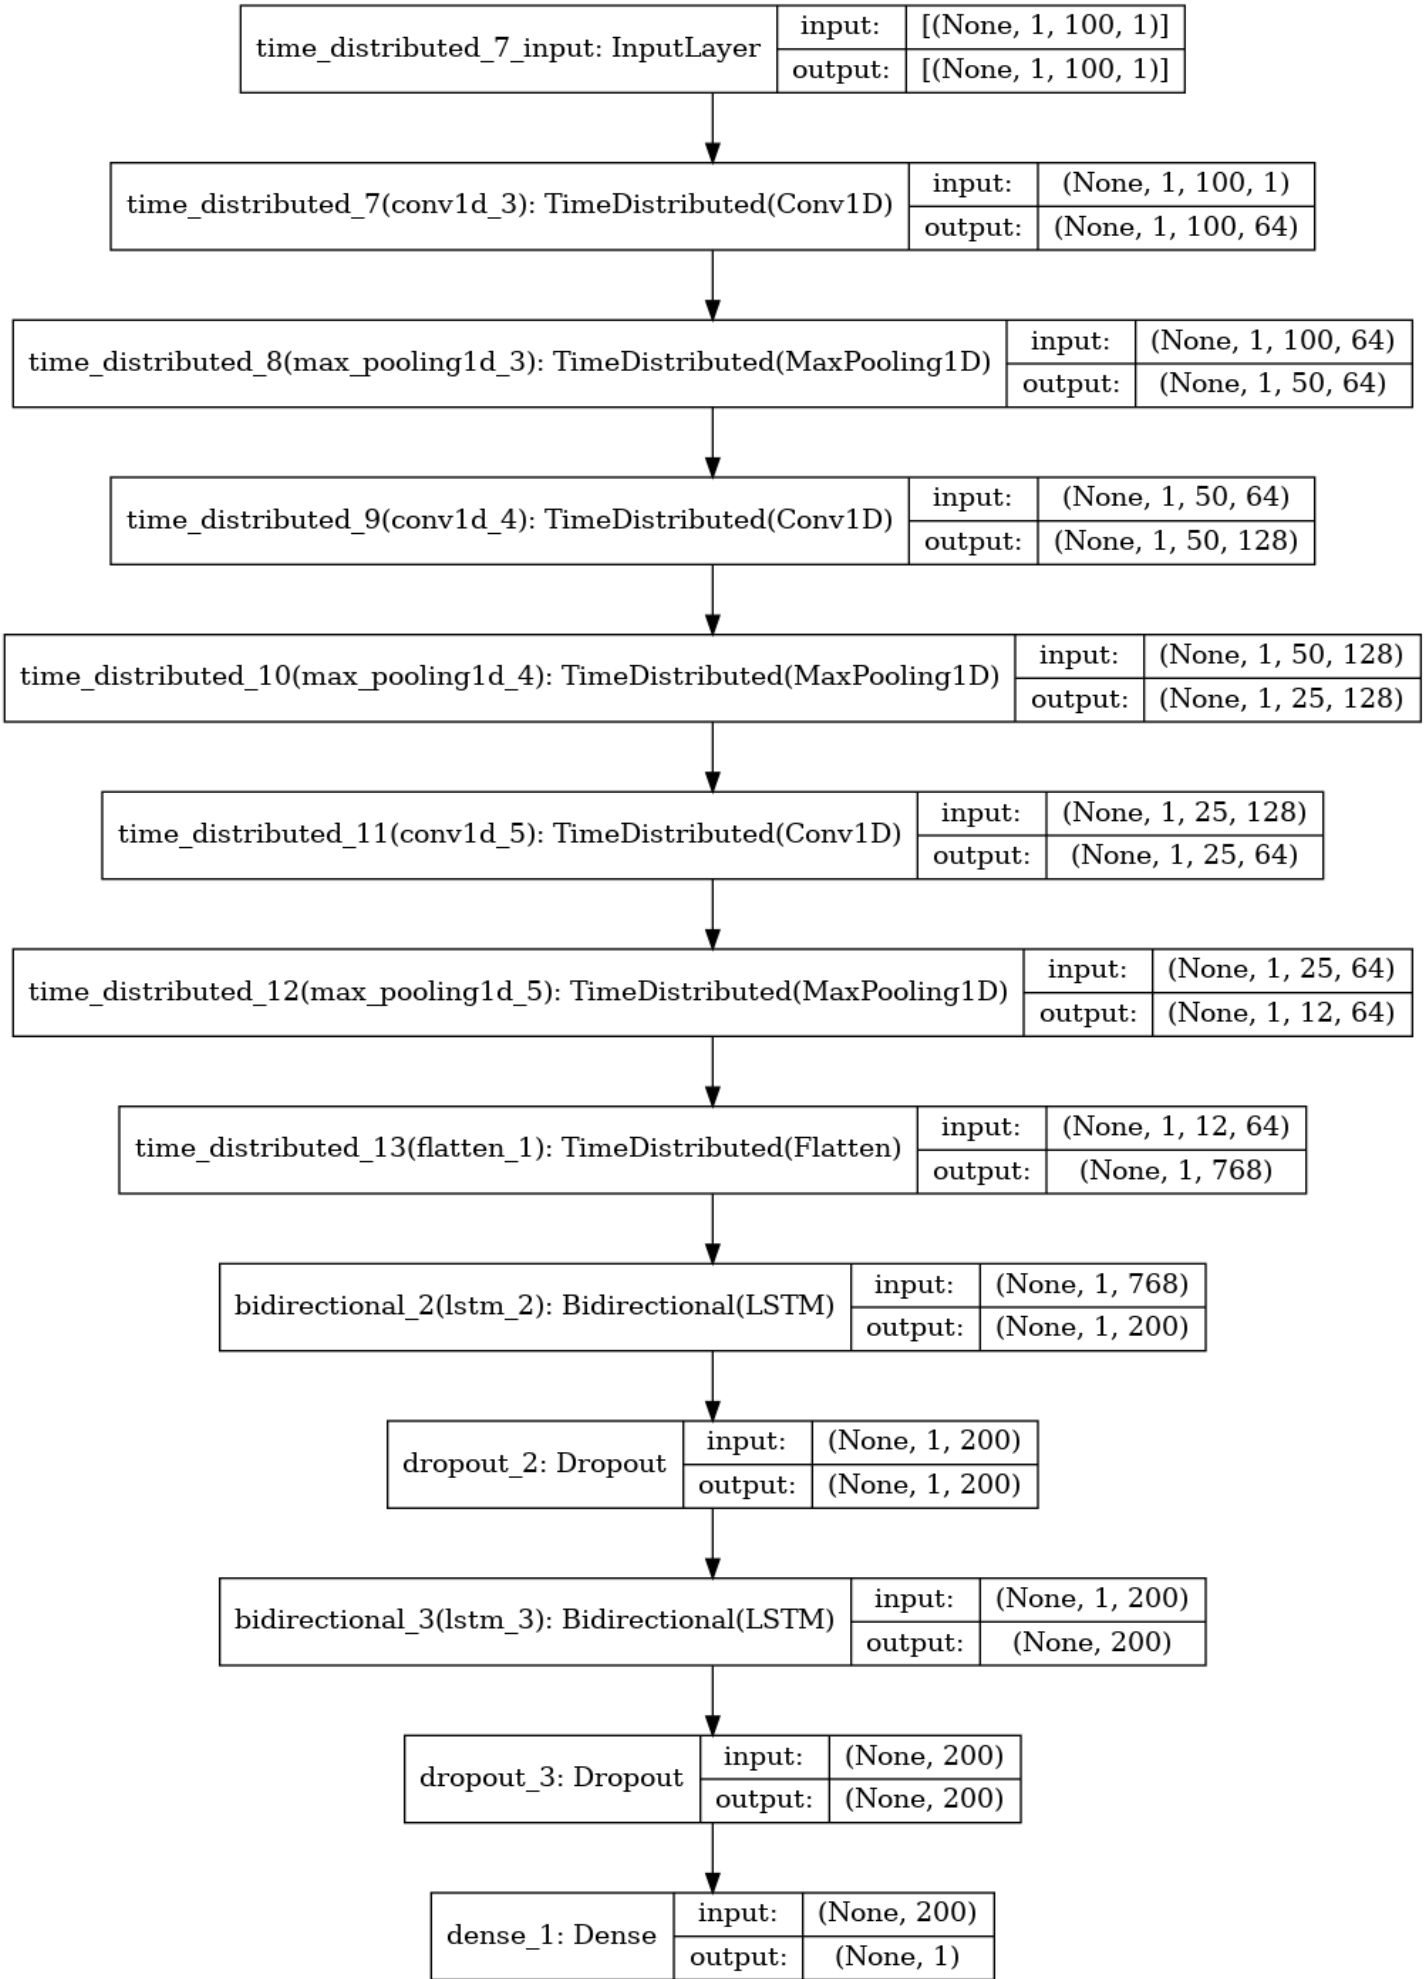
Model: "sequential_1"

Layer (type)	Output Shape	Param #
time_distributed_7 (TimeDist (None, 1, 100, 64))		128
time_distributed_8 (TimeDist (None, 1, 50, 64))		0
time_distributed_9 (TimeDist (None, 1, 50, 128))		8320
time_distributed_10 (TimeDis (None, 1, 25, 128))		0
time_distributed_11 (TimeDis (None, 1, 25, 64))		8256
time_distributed_12 (TimeDis (None, 1, 12, 64))		0
time_distributed_13 (TimeDis (None, 1, 768))		0
bidirectional_2 (Bidirection (None, 1, 200))		695200
dropout_2 (Dropout)	(None, 1, 200)	0
bidirectional_3 (Bidirection (None, 200))		240800
dropout_3 (Dropout)	(None, 200)	0
dense_1 (Dense)	(None, 1)	201
Total params: 952,905		

Trainable params: 952,905
Non-trainable params: 0

None

Out[60]:



In [61]:

```
model.evaluate(test_X, test_Y)
```

```
20/20 [=====] - 0s 4ms/step - loss: 0.0015 - mse: 0.0015 - mae: 0.0277
```

Out[61]:

```
[0.0014527516905218363, 0.0014527516905218363, 0.027694158256053925]
```

In [62]:

```
from sklearn.metrics import explained_variance_score
from sklearn.metrics import r2_score
from sklearn.metrics import max_error

# predict probabilities for test set
yhat_probs = model.predict(test_X, verbose=0)
# predict crisp classes for test set
yhat_classes = model.predict_classes(test_X, verbose=0)
# reduce to 1d array
yhat_probs = yhat_probs[:, 0]
yhat_classes = yhat_classes[:, 0]

var = explained_variance_score(test_Y.reshape(-1,1), yhat_probs)
print('Variance: %f' % var)

r2 = r2_score(test_Y.reshape(-1,1), yhat_probs)
print('R2 Score: %f' % var)

var2 = max_error(test_Y.reshape(-1,1), yhat_probs)
print('Max Error: %f' % var2)
```


```
Variance: 0.947681
R2 Score: 0.947681
Max Error: 0.216366
```

```
/opt/conda/lib/python3.7/site-packages/tensorflow/python/keras/engine/sequential.py:450:
UserWarning: `model.predict_classes()` is deprecated and will be removed after 2021-01-01
. Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your model does multi-
class classification (e.g. if it uses a `softmax` last-layer activation). * `(model.pre
dict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it
uses a `sigmoid` last-layer activation).
warnings.warn("`model.predict_classes()` is deprecated and '
```

In [63]:

```
predicted = model.predict(test_X)
test_label = test_Y.reshape(-1,1)
predicted = np.array(predicted[:,0]).reshape(-1,1)
len_t = len(train_X)
for j in range(len_t, len_t + len(test_X)):
    temp = data.iloc[j,3]
    test_label[j - len_t] = test_label[j - len_t] * temp + temp
    predicted[j - len_t] = predicted[j - len_t] * temp + temp
plt.plot(predicted, color = 'green', label = 'Predicted Stock Price')
plt.plot(test_label, color = 'red', label = 'Real Stock Price')
plt.title(' Stock Price Prediction')
plt.xlabel('Time')
plt.ylabel(' Stock Price')
plt.legend()
plt.show()
```





Testing part

In this part, the model is saved and loaded back again. Then, it's made to train again but with different data to check it's loss and prediction

In [64]:

```
# First we need to save a model
model.save("model.h5")
```

In [65]:

```
# Load model
new_model = tf.keras.models.load_model("./model.h5")
```

In [66]:

```
new_model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
time_distributed_7 (TimeDist	(None, 1, 100, 64)	128
time_distributed_8 (TimeDist	(None, 1, 50, 64)	0
time_distributed_9 (TimeDist	(None, 1, 50, 128)	8320
time_distributed_10 (TimeDis	(None, 1, 25, 128)	0
time_distributed_11 (TimeDis	(None, 1, 25, 64)	8256
time_distributed_12 (TimeDis	(None, 1, 12, 64)	0
time_distributed_13 (TimeDis	(None, 1, 768)	0
bidirectional_2 (Bidirection	(None, 1, 200)	695200
dropout_2 (Dropout)	(None, 1, 200)	0
bidirectional_3 (Bidirection	(None, 200)	240800
dropout_3 (Dropout)	(None, 200)	0
dense_1 (Dense)	(None, 1)	201
Total params: 952,905		
Trainable params: 952,905		
Non-trainable params: 0		

In [67]:

```
# For data preprocessing and analysis part
#data2 = pd.read_csv('../input/price-volume-data-for-all-us-stocks-etfs/Stocks/aaio.us.tx
t')
#data2 = pd.read_csv('../input/nifty50-stock-market-data/SBIN.csv')
#data2 = pd.read_csv('../input/stock-market-data/stock_market_data/nasdaq/csv/ACTG.csv')
```



```

data2 = pd.read_csv('./data.csv')
# Any CSV or TXT file can be added here....
data2.dropna(inplace=True)
data2.head()

data2.reset_index(drop=True, inplace=True)
data2.fillna(data.mean(), inplace=True)
data2.head()
df2 = data2.drop('date', axis=1)

print(df2)

X = []
Y = []
window_size=100
for i in range(1, len(df2) - window_size - 1, 1):
    first = df2.iloc[i,4]
    temp = []
    temp2 = []
    for j in range(window_size):
        temp.append((df2.iloc[i + j, 4] - first) / first)
    # for j in range(week):
    temp2.append((df2.iloc[i + window_size, 4] - first) / first)
    # X.append(np.array(stock.iloc[i:i+window_size,4]).reshape(50,1))
    # Y.append(np.array(stock.iloc[i+window_size,4]).reshape(1,1))
    # print(stock2.iloc[i:i+window_size,4])
    X.append(np.array(temp).reshape(100, 1))
    Y.append(np.array(temp2).reshape(1, 1))

x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, shuffle=True)

train_X = np.array(x_train)
test_X = np.array(x_test)
train_Y = np.array(y_train)
test_Y = np.array(y_test)

train_X = train_X.reshape(train_X.shape[0],1,100,1)
test_X = test_X.reshape(test_X.shape[0],1,100,1)

print(len(train_X))
print(len(test_X))

```

	Unnamed: 0	open	high	low	close	adjusted close	volume \
0	0	115.00	116.335	114.56	115.81	115.810000	3322012
1	1	116.16	117.270	116.08	116.73	116.730000	3220802
2	2	116.79	117.940	116.04	116.79	116.790000	4914995
3	3	116.00	118.810	115.19	116.47	116.470000	6417218
4	4	116.49	116.560	115.27	116.05	116.050000	5384548
...
5550	5550	92.75	92.940	90.19	90.25	52.226608	13737600
5551	5551	94.44	94.440	90.00	91.56	52.984689	16697600
5552	5552	95.87	95.940	93.50	94.37	54.610803	10369100
5553	5553	96.75	96.810	93.69	94.81	54.865426	11105400
5554	5554	98.50	98.810	96.37	96.75	55.988081	9551800

	dividend amount	split cf
0	0.0	1.0
1	0.0	1.0
2	0.0	1.0
3	0.0	1.0
4	0.0	1.0
...
5550	0.0	1.0
5551	0.0	1.0
5552	0.0	1.0
5553	0.0	1.0
5554	0.0	1.0


```

[5555 rows x 9 columns]
4362
1091

```

In [68]:

```
predicted = model.predict(test_X)
test_label = test_Y.reshape(-1,1)
predicted = np.array(predicted[:,0]).reshape(-1,1)
len_t = len(train_X)
for j in range(len_t, len_t + len(test_X)):
    temp = data2.iloc[j,3]
    test_label[j - len_t] = test_label[j - len_t] * temp + temp
    predicted[j - len_t] = predicted[j - len_t] * temp + temp
plt.plot(predicted, color = 'green', label = 'Predicted Stock Price')
plt.plot(test_label, color = 'red', label = 'Real Stock Price')
plt.title(' Stock Price Prediction')
plt.xlabel('Time')
plt.ylabel(' Stock Price')
plt.legend()
plt.show()
```



In []: